

- [MangoDB Technical Report](#)
  - [Table of Contents](#)
  - [System Architecture Overview](#)
  - [Core Components](#)
    - [1. Document Storage](#)
    - [2. Transaction Management](#)
      - [ACID Properties Implementation](#)
  - [Concurrency Control](#)
    - [1. Locking Mechanism](#)
    - [2. Deadlock Detection](#)
  - [Data Storage and Persistence](#)
    - [1. File Structure](#)
    - [2. Indexing System](#)
  - [Query Processing](#)
    - [1. Query Pipeline](#)
    - [2. Query Execution](#)
  - [Recovery System](#)
    - [1. Checkpointing](#)
    - [2. Log-Based Recovery](#)
  - [Limitations and Trade-offs](#)
    - [Current Limitations](#)
    - [Design Decisions and Trade-offs](#)
  - [Future Improvements](#)

# MangoDB Technical Report

---

A detailed technical analysis of our document-based database system implementation.

## Table of Contents

---

- [1. System Architecture Overview](#)
- [2. Core Components](#)
- [3. Transaction Management](#)
- [4. Concurrency Control](#)
- [5. Data Storage and Persistence](#)

6. [Query Processing](#)
7. [Recovery System](#)
8. [Limitations and Trade-offs](#)

# System Architecture Overview

---

MangoDB is built as a layered architecture:

Web Interface	Flask-based UI
Query Parser	MongoDB-like syntax
Transaction Mgr	ACID & Concurrency
Storage Layer	JSON files & Indexes

## Core Components

---

### 1. Document Storage

**Concept:** Document-based storage allows flexible schema and nested data structures.

**Implementation:**

```
class DocumentDB:
    def __init__(self):
        self.databases_dir = "databases/"
        self._ensure_databases_dir()
```

Each document is stored as a JSON object with a unique UUID:

```
doc_id = str(uuid.uuid4())
document = {
    "_id": doc_id,
    "data": {...}
}
```

## 2. Transaction Management

### ACID Properties Implementation

#### Atomicity:

- Each transaction is wrapped in a try-except block
- Operations are logged before execution
- Two-phase commit protocol:
  1. Prepare phase: Log operations
  2. Commit phase: Apply changes

```
def execute_transaction(self, transaction_id):
    try:
        # Log operation
        self.log_operation(transaction_id, operation)
        # Execute operation
        result = self.execute_operation(operation)
        # Commit if successful
        self.commit_transaction(transaction_id)
        return result
    except Exception:
        self.abort_transaction(transaction_id)
        raise
```

#### Consistency:

- Document validation before writes
- Schema enforcement through JSON structure
- Referential integrity checks

#### Isolation Levels:

```
class IsolationLevel(Enum):
    READ_UNCOMMITTED = 1
    READ_COMMITTED = 2
    REPEATABLE_READ = 3
    SERIALIZABLE = 4
```

Implementation in `transaction_manager.py`:

```
def begin_transaction(self, isolation_level):
    transaction_id = str(uuid.uuid4())
    self.transactions[transaction_id] = {
        "state": TransactionState.ACTIVE,
        "isolation_level": isolation_level,
        "locks": set()
    }
    return transaction_id
```

### Durability:

- Transaction logging
- Periodic checkpointing
- Crash recovery mechanism

# Concurrency Control

## 1. Locking Mechanism

**Concept:** Document-level locking prevents concurrent modifications to the same document.

### Implementation Flow:

```
Transaction Start
  ↓
Request Lock
  ↓
Check Lock Table
  ↓
Grant/Block Lock
  ↓
Operation Execution
  ↓
Release Lock
```

### Lock Types:

```
class LockType(Enum):
    READ = 1
    WRITE = 2
```

## Lock Manager Implementation:

```
class LockManager:
    def __init__(self):
        self.lock_table = {} # doc_id -> {lock_type, transaction_id}

    def acquire_lock(self, doc_id, lock_type, transaction_id):
        if self._is_lock_compatible(doc_id, lock_type):
            self.lock_table[doc_id] = {
                "type": lock_type,
                "transaction_id": transaction_id
            }
            return True
        return False
```

## 2. Deadlock Detection

### Implementation:

- Timeout-based approach
- Lock wait graph construction
- Cycle detection

```
def detect_deadlock(self, transaction_id):
    wait_graph = self._build_wait_graph()
    if self._has_cycle(wait_graph):
        self._resolve_deadlock(transaction_id)
```

## Data Storage and Persistence

### 1. File Structure

```
databases/
├── db_name/
│   ├── collection.json
│   └── indexes/
│       └── field_name.idx
├── transaction_logs/
└── checkpoints/
```

## 2. Indexing System

### B+ Tree Implementation:

```
class BPlusTree:
    def __init__(self):
        self.root = None
        self.order = 4 # Maximum children per node

    def insert(self, key, value):
        if not self.root:
            self.root = LeafNode()
        return self._insert_recursive(self.root, key, value)
```

## Query Processing

### 1. Query Pipeline

```
Query String
  ↓
Parser (query_parser.py)
  ↓
Optimizer (queryOptimizer.py)
  ↓
Executor (queryProcessor.py)
  ↓
Result
```

## 2. Query Execution

```
def execute_query(self, query):
    # Parse query
    operation, collection, params = parse_raw_query(query)

    # Start transaction
    transaction_id = self.begin_transaction()

    try:
        # Execute operation
        result = self._execute_operation(operation, collection, params)
    # Commit
```

```
        self.commit_transaction(transaction_id)
    return result
except Exception:
    self.abort_transaction(transaction_id)
    raise
```

# Recovery System

---

## 1. Checkpointing

- Every 60 seconds
- Saves system state
- Maintains last 5 checkpoints

## 2. Log-Based Recovery

```
def recover_from_crash(self):
    # Load latest checkpoint
    checkpoint = self._load_latest_checkpoint()

    # Replay logs
    logs = self._get_logs_after_checkpoint(checkpoint)
    for log in logs:
        if log["status"] == "committed":
            self._redo_operation(log)
        else:
            self._undo_operation(log)
```

# Limitations and Trade-offs

---

## Current Limitations

### 1. Scalability:

- File-based storage limits concurrent access
- No distributed architecture
- Memory constraints for large datasets

## 2. Performance:

- JSON parsing overhead
- File I/O bottlenecks
- Limited index types

## 3. Features:

- No complex queries
- Limited aggregation
- Basic indexing only

# Design Decisions and Trade-offs

## 1. Document-Level vs Collection-Level Locking

- Chose document-level for better concurrency
- Trade-off: More complex lock management

## 2. File-Based vs In-Memory Storage

- Chose file-based for durability
- Trade-off: Slower performance

## 3. B+ Tree Indexing

- Chose B+ tree for range queries
- Trade-off: More complex implementation

# Future Improvements

---

## 1. Performance:

- Implement caching
- Add more index types
- Optimize file I/O

## 2. Features:

- Add complex queries



- Implement aggregation
- Support more data types

### 3. **Scalability:**

- Add sharding
- Implement replication
- Support distributed architecture